

LAB3, Polling and Interrupt-Driven I/O

Goal

Become aware of the various ways to communicate with the peripherals. Understand the advantages and disadvantages of each.

Given files

Assuming you copied the lab files according to the lab setup instruction you can find the files you need in `~/TDDI11/lab3`. Your modifications goes to `serial.asm`.

Assignment

You will write a program that reads characters typed on the keyboard and sends the read characters on the serial interface. At the same time, it displays characters read from the serial interface.

Write assembly code to write the serial port by *polling*. With polling you continuously test a bit of a register. If the bit is not set, the serial interface is busy sending another character, and therefore you should not write the new character to the serial interface. If the bit is set, the serial interface is accepting new characters to be send.

Write the assembly code to read characters from the serial interface by *interrupts*. Upon receiving an interrupt, the serial interface controller notifies the CPU by means of an interrupt. The serial port share interrupt signal with other devices, so you must still check that a byte really arrived at the serial port before reading it.

Edit the file `serial.asm` to fill-in the missing details of the polled waiting-loop output function (`SerialPut`) and the serial input ISR (`SerialISR`).

In order to be able to test the program, you need to emulate two computers connected by a serial line. Just issue the command `gmake link` in your lab 3 folder to do just that.

Deliverables

Code of `SerialPut` and `SerialISR`. Demo for the lab assistant.

Background

We discuss two ways of communicating with peripherals: polling and interrupts.

Polling means continuously checking the peripheral to detect if it has changed state. For example, if the keyboard is expected to place the ASCII code of a pressed key in a certain register of the keyboard controller, polling would imply continuously reading that register to see if the value changes. We can see immediately the disadvantage: instead of using the processor to do something useful, we waste time checking something that happens with a low rate.

An alternative to polling is interrupt-driven communication. Instead of the processor interrogating the peripheral, the peripheral notifies the processor if its state has changed. The peripheral asserts a signal that interrupts the execution of the processor. The execution jumps to a predefined address where the Interrupt Service Routine (ISR) resides. The ISR implements the response to the interrupt. In this way, the processor works with the peripheral only when needed.

Nevertheless, polling may have advantages too. For example, there might be intervals when we want to ignore the peripheral and to run uninterrupted. If we used interrupt-driven communication the processor would be interrupted even if it decides not to service the interrupt. Temporarily disabling the interrupts may not be a good idea if we want not to be interrupted, because we could be interested in other interrupts. For example, we want to ignore the keyboard but not the drop of the battery current below a threshold. In this case, we would communicate with the keyboard by polling, i.e. we consult it only when we want that.

Polling can also have advantage when communicating with a device that operate almost as fast as the processor. If we use interrupts many cycles will go to waste in the overhead of doing the interrupt, while polling may be able to read one data every iteration in the poll loop.

In such cases DMA is an other solution, in which a large block of data is transferred to memory without CPU intervention, and a interrupt signal delivered when the transfer is done, and the CPU need to provide a new empty buffer. While the new buffer is filled, the CPU can process the previous buffer content in parallel.

The IBM-PC Serial port

The IBM-PC typically supports two (or more) serial ports, called COM1 and COM2. These devices are implemented using an I/O chip called a Universal Asynchronous Receiver-Transmitter (UART). UARTs convert between the 8-bit parallel format used inside the PC and a serial representation in which information is transmitted one bit at a time.

Each serial port is assigned a block of eight I/O port addresses, starting with a “base” address of 0x2f8, 0x3f8, 0x2e8, or 0x3e8. These addresses typically correspond to COM1, COM2, COM3, and COM4 respectively. These port addresses are used with I/O instructions to write data or commands into the internal registers of the UART, or to read data or status information back out.

The three most important registers are:

1) LSR (Line Status Register) I/O Port: Base+5 (Read-Only)

7	6	5	4	3	2	1	0
FIFO	Transmitter	THRE	Break	Framing	Parity	Overrun	RBF
Error	Empty		Detected	Error	Error	Error	

For the purpose of the labs, only two status bits are important:

- Bit 5: THRE (Transmitter Holding Register Empty): This bit is 1 when the UART is ready to transmit another byte of data out on the serial line. Wait for this bit to become 1 before writing to THR.
 - Bit 0: RBF (Receiver Buffer Full): If this bit is 1, input data is available in RBR.
- 2) THR (Transmitter Holding Register) I/O Port: Base+0 (Write-Only) Data to be transmitted serially to a remote computer may be written to this port when bit THRE in the LSR is 1.
 - 3) RBR (Receiver Buffer Register) I/O Port: Base+0 (Read-Only) Data received serially from a remote computer may be read from this port when bit RBF in the LSR is 1.

To write a byte to the serial port you must first make sure the device is ready. This is done by reading the Line Status Register and extract bit 5, Transmitter Holding Register Empty. If it is set it means the Transmitter Holding Register is empty, and you can safely proceed by writing to the Transmitter Holding Register located at the base port. If bit 5 is not set you must first wait for it to become set. (poll it).

To read a byte from the serial port by using interrupts the address to an interrupt handler must be places in the processors interrupt service vector. (In the lab this is already done for you.) In the interrupt handler you must then determine why the processor was interrupted. Besides getting a character on the serial line, other events such as “break detected” or all kind of errors may cause interrupts. To see if a new character is available on the serial port read the Line Status Register and extract bit 0, Receive Buffer Full. If it is set a character has arrived on the line, and you can safely proceed by fetching it from the base port. If it is not set it means the interrupt was for something else and you should do nothing (as we only care about the serial port). In both cases you have to send the End Of Interrupt (EOI) to the interrupt controller. This is done by writing 0x20 to port 0x20.

The x86 I/O Instructions

The 80x86 supports two I/O instructions: `in` and `out`. They are like very limited versions of the `mov` instruction, with the major difference that they do not access the normal memory bus, but instead a special I/O bus, with 65536 I/O ports (a port is an address on the I/O bus, but named port to distinguish it from a memory bus address). The I/O instructions take the forms:

```
in      eax/ax/al, port
in      eax/ax/al, dx
out     port,  eax/ax/al
out     dx,   eax/ax/al
```

The `in` instruction must always have (part of) `eax` as destination, and can read the port address either as a directly specified 8-bit port number (0 – 0xFF) or use the full 16-bit port number in `edx`. In most cases you must thus load the port number to `edx` first.

The `out` instruction is exactly the same, but reverses the source and destination.

Neither instruction modifies any flags in the flags register.

Resources

To better understand I/O with polling and interrupts the structure of the x86 system you can read the relevant chapters in:

<http://flint.cs.yale.edu/cs422/doc/art-of-asm/pdf/>

Relevant chapters:

CH03-6

CH17-3

CH22-1

