

# x86 Assembly

---

CS350 – Undergraduate Operating Systems  
02/10/2021

Sasan Golchin



# Today's agenda

---

- Hints about the Primer
- x86 Assembly
  - Architecture Overview
  - The AT&T Operand Format
  - Data Movement Instructions
  - Arithmetic Instructions
  - Bitwise Logical Instructions
  - Control Flow Instructions
  - GNU Inline Assembly



# How to go about the Keyboard Driver?

---

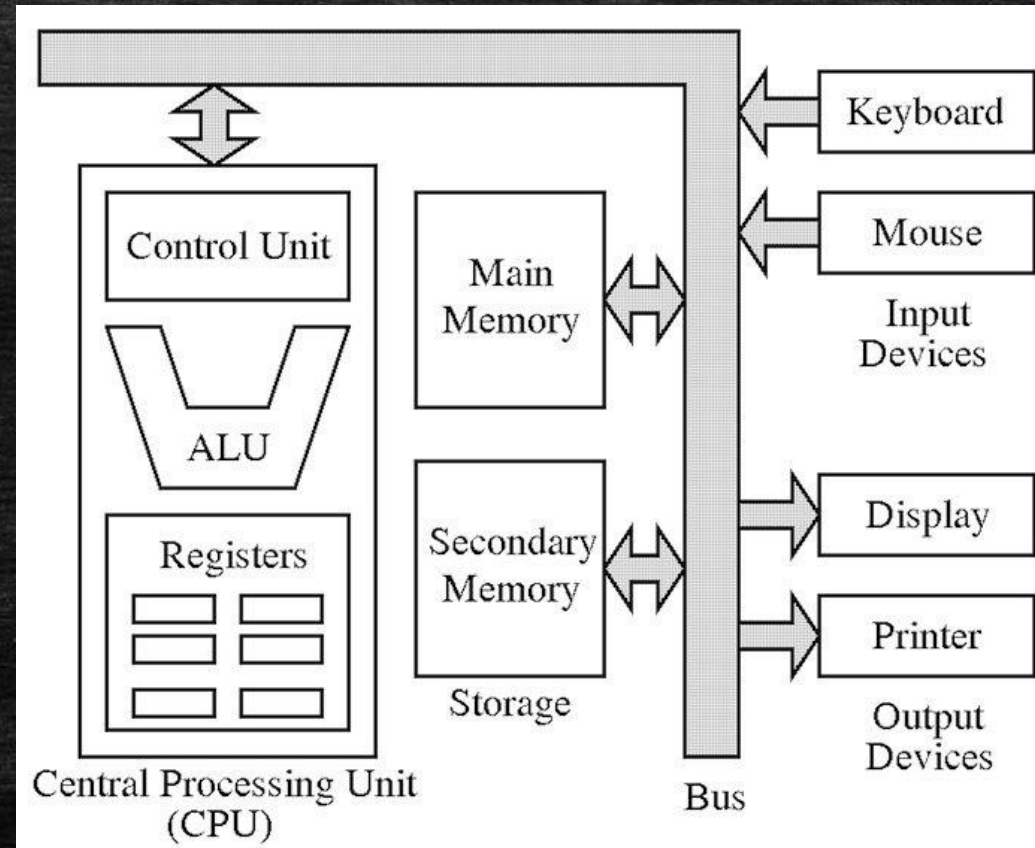
- Start as early as you can!
- Start by extending the provided IOCTL module/test code
- Polling method: You need to disable the native driver (i8042).
- Interrupt Based: Look for Linux kernel functions to define and install your own ISR. Google it!
- At this point, you will get double characters as you are co-existing with the native i8042! Look for a way to bypass that.
- Distinguish key-press from key-release events. Look at the bit definition of the data port (0x60). Read [this](#) and then [this](#)!
- Work out how to manipulate wait queues in Linux. You need to make the caller task block until there's a keyboard event!



# x86 Architecture Overview

## Von Neumann Architecture

- Unified memory for Instructions and Data, i.e., the Main Memory
- Move data between Main Memory and Registers
- System Bus:
  - Data Bus: 32-bit wide
  - Address Bus: 32-bit wide
  - Control Bus: R/W, Memory/IO





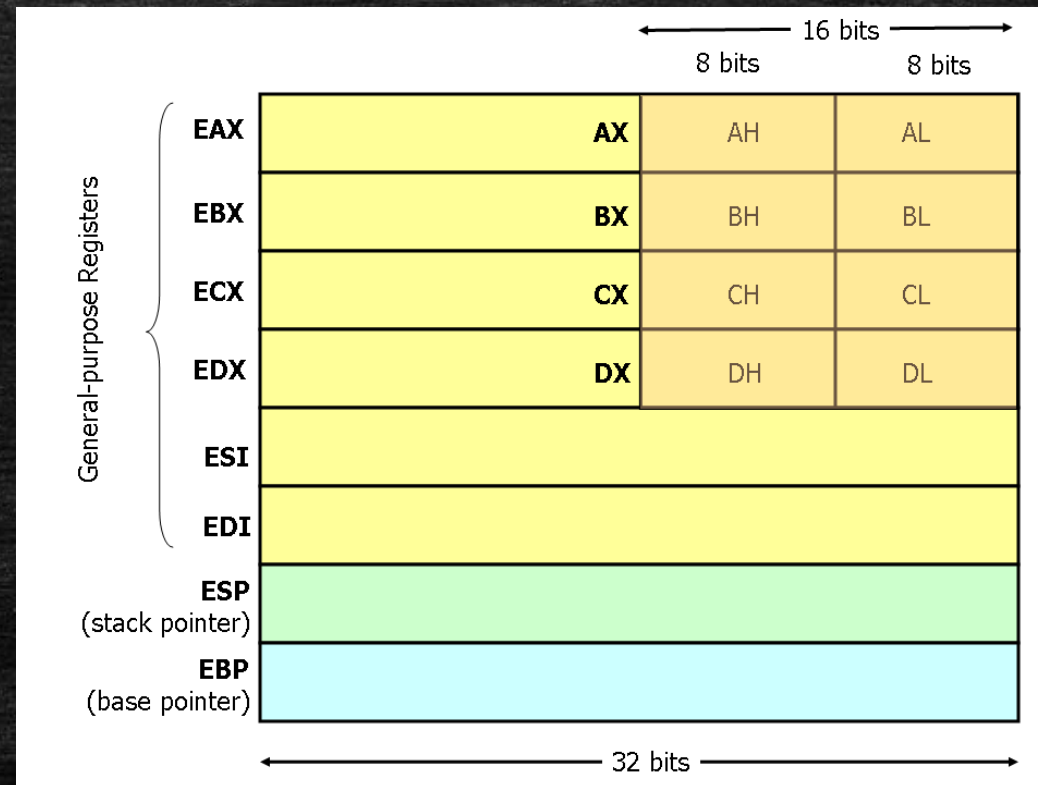
# x86 Registers

## General Purpose Registers

- EAX, EBX, ECX, EDX
- AX, BX, CX, DX
- AH, AL, BH, BL, ...
- ESI, EDI

## Special Purpose Register

- ESP: Stack pointer
- EBP: Base pointer
- EIP: Instruction Pointer
- And a lot more: Control Registers, Model-Specific Registers and so on....





# Program Organization (1)

---

- A program is a collection of Data and Code spread over different sections.
- We use GNU-AS to compile our x86 assembly code.
- Some useful GNU-AS directives to organizing your program:
  - **Definition of sections:**
    - .data, .rodata, .bss, .text, .section [name-of-custom-section]
  - **Definition of labels:** Labels are name of constants, variables, functions and anything that can be addressed in the program!
    - [name-of-label]:
  - **Data definition directives:**
    - .byte, .short, .long, .zero, .string, .space, .float, .double



# Program Organization (2)

---

- Examples data section of an assembly program

```
.data /* initialized global variables */
my_byte_arr: /* define an array of bytes labeled my_byte_arr. */
    .byte 64, 0x10, 0xFF
x: /* define a 2-byte integer variable labeled x initialized to 42 */
    .short 42
y: /* define a 4-byte integer variable labeled y = 0x1234ABCD */
    .long 0x1234ABCD
s1: /* define a null-terminated string initialized to "Hello World" */
    .string "Hello World"

.bss /* Uninitialized global variables */
buf: /* Reserve 256 bytes in a buffer labeled buf */
    .space 256

.text
/* This is where our code (x86 instructions) goes! */
```



# An Introduction to x86 Instruction Set

---

- Data Movement Instructions
  - mov, push, pop, lea, in, out
- Arithmetic Instructions
  - add, sub, inc, dec, imul, idiv
- Logical Instructions
  - and, or, xor, not, neg, shl, shr, sar
- Control-Flow (Branching) Instructions
  - jmp, je, jne, jz, jnz, jg, jge, jl, jle
  - call, ret
  - int
- Many more that we can't possibly cover in this lab.



# The AT&T Syntax - Operands

- The AT&T syntax for instructions w/ more than one operand:  
**INSTR src, dst** or **INSTR src1, src2, dst**
- Let's consider the MOV instruction. It is used to move data between the registers and the main memory.
- The source and destination operands can be one of the following:
  1. **Registers:** %[register name] e.g., mov %eax, %ebx
  2. **Immediate:** \$[constant value] e.g., mov \$0x10, %eax
  3. **Memory Location:** Follows the format **Offset(Base, Index, Scale)**
    - **Offset** is an immediate (Constant number or a label)
    - **Base** and **Index** are x86 registers
    - **Scale** is an immediate (Constant number or a label)
    - Memory Location = Offset + (Base + Index\*Scale)



# The AT&T Syntax - Operands

---

- Sometimes the instruction operands cannot unequivocally specify how many bytes should the instruction operate on. E.g.,:
- Copying some content from address X in the memory into EAX: Should the CPU copy 1, 2 or 4 bytes from X?
- The following suffixes are attached to the name of the instruction to clarify the size of the operand:
  - **b**: One byte
  - **w**: A word, i.e., Two bytes
  - **l**: A long word, i.e., Four bytes
- Then the instruction mov has four forms: mov, movb, movw, movl



# Data Movement Instructions (1)

---

## Move!

- Syntax

- `mov <reg>, <reg>`
- `mov <reg>, <mem>`
- `mov <mem>, <reg>`
- `mov <imm>, <reg>`
- `mov <imm>, <mem>`

- Remember: The first operand is the source and the last one is the destination.
- Also remember that the `<mem>` operands follow the format: **Offset**(Base, Index, Scale)



# MOV Examples

1. **Direct Memory:** `MOVb var(,1), %ecx`

- Write the 1-byte value at memory address of the label var into ECX. The suffix b means "a byte"! Label are translated into address values by the linker.

2. **Indirect Memory:** `MOVw (%ebx), %eax`

- Write the 2-byte value at memory address stored in EBX into EAX – i.e., Dereference EBX into EAX. The suffix w means "a word" that is 2 bytes.

3. **Indexed Memory:** `MOVL 8(%ebx,%esi,4), %edx`

- Move the 4-bytes value at address  $8 + (EBX + ESI * 4)$  into EDX. The suffix l means "a long word" that is 4 bytes
- Assume you want to address the 5<sup>th</sup> element of an array of 4-byte integers located at offset 8 of a data-structure whose base address is at 0x100000

```
mov $0x100000, %ebx      /* Set the base register */
mov $5, %esi             /* Set the index register */
movl 8(%ebx, %esi, 4), %edx /* access the main memory */
```



# Data Movement Instructions (2)

---

LEA (Load Effective Address): Computes the absolute value of a memory location specified in the **Offset(Base, Index, Scale)** format.

Think of MOV as dereferencing a pointer and LEA as reading the address in a pointer.

## Syntax

- lea <mem>, <reg32>

## ▪ Examples

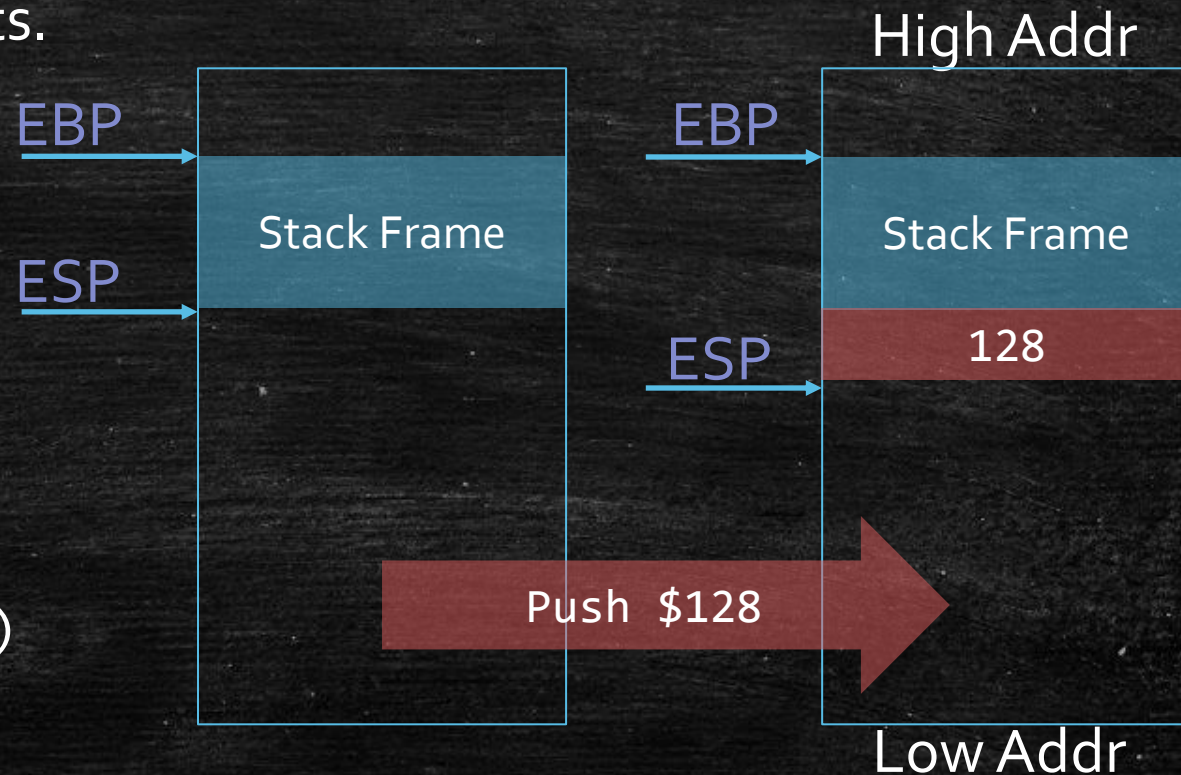
- lea (%ebx,%esi,8), %edi /\* EDI <- EBX+8\*ESI \*/
- lea val(,1), %eax /\* EAX <- val \*/



# Data Movement Instructions (3)

PUSH: Places its operand onto the top of the hardware supported stack in memory, where ESP points.

- Syntax
  - push <reg>
  - push <mem>
  - push <imm>
- It's Equivalent to:
  - Decrement ESP by 4
  - movl <operand>, (%esp)



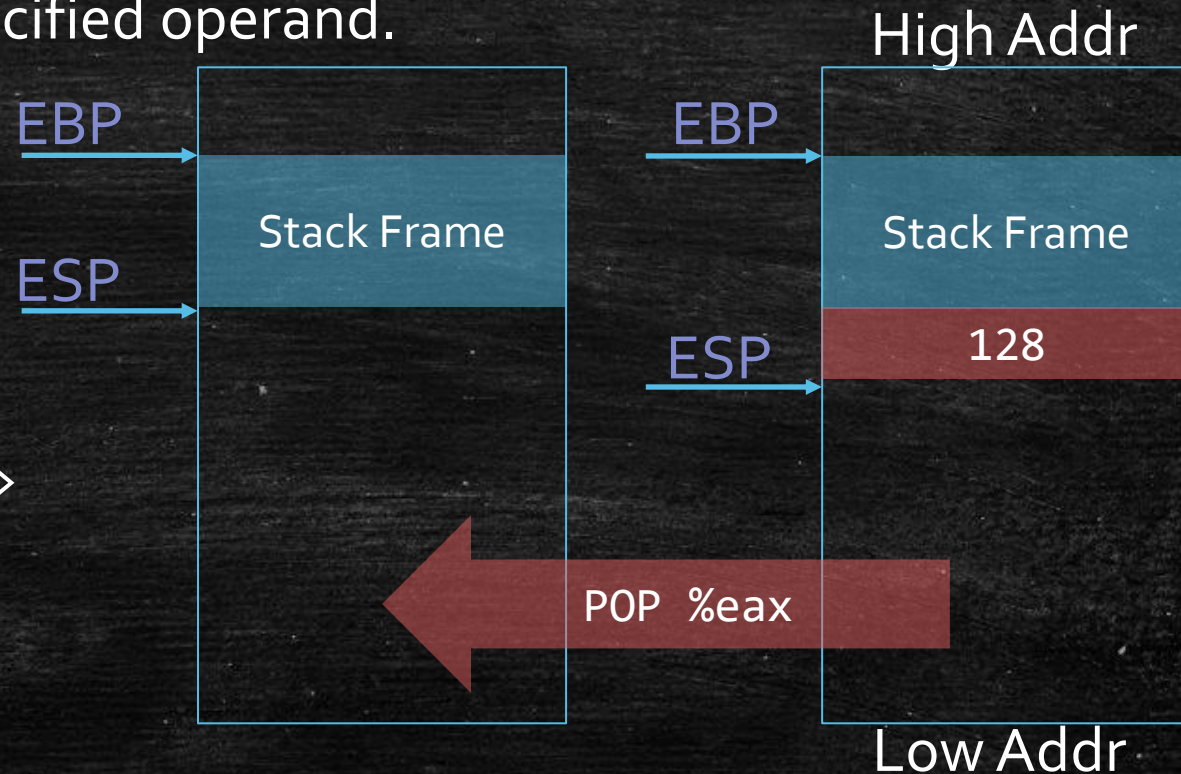


# Data Movement Instructions (4)

POP: Removes the 4-byte data element from the top of the hardware-supported stack into the specified operand.

## Syntax

- push <reg>
- push <mem>
- It's Equivalent to:
  - movl (%esp), <operand>
  - Increase ESP by 4





# Data Movement Instructions (5)

---

## Reading and Write Hardware I/O Ports

- Syntax

- `in <imm8>, <AL, AX, or EAX>`
- `in DX, <AL, AX, or EAX>`
- `out <AL, AX, or EAX>, <imm8>`
- `out <AL, AX, or EAX>, DX`

- Examples

- `inb $0x64, %al /* Read one byte from port# 0x64 into AL */`
- `inw %dx, %ax /* Read two bytes from port# in DX into AX */`
- `outb %al, %dx /* Write the byte in AL into the port in DX */`



# Arithmetic Instructions (1)

---

## Integer Addition

- Syntax

- add <reg>, <reg>
- add <mem>, <reg>
- add <reg>, <mem>
- add <imm>, <reg>
- add <imm>, <mem>

- Examples

- add \$10, %eax /\* EAX is set to EAX + 10 \*/
- addb \$10, (%eax) /\* add 10 to the single byte stored at memory address stored in EAX \*/



# Arithmetic Instructions (2)

---

## Integer Subtraction

- Syntax

- sub <reg>, <reg>
- sub <mem>, <reg>
- sub <reg>, <mem>
- sub <imm>, <reg>
- sub <imm>, <mem>

- Examples

- sub %ah, %al /\* AL <- AL - AH \*/
- sub \$55, %eax /\* EAX <- EAX - 55 \*/



# Arithmetic Instructions (3)

---

## Increment, Decrement

- Syntax

- `inc <reg>`
- `inc <mem>`
- `dec <reg>`
- `dec <mem>`

- Examples

- `dec %eax`      `/* subtract 1 from the contents of EAX */`
- `incl var(,1)` `/* add 1 to the 32-bit int. at location var */`



# Arithmetic Instructions (4)

---

## Integer Multiplication

- Syntax

- `imul <reg32>, <reg32>`
- `imul <mem>, <reg32>`
- `imul <con>, <reg32>, <reg32>`
- `imul <con>, <mem>, <reg32>`

- Examples

- `imul (%ebx), %eax /* EAX <- EAX * 32-bit value @ mem[EBX] */`
- `imul $25, %edi, %esi /* ESI <- EDI * 25 */`



# Arithmetic Instructions (5)

---

## Integer Division

- Syntax
  - `idiv <reg32>`
  - `idiv <mem>`
- Divides the contents of the 64-bit integer EDX:EAX by the specified operand value. The quotient result of the division is stored into EAX, while the remainder is placed in EDX.
- Examples
  - `idiv %ebx`     `/* EDX:EAX / EBX */`
  - `idivw (%ebx)` `/* EDX:EAX / <16-bit value at mem[EBX]> */`



# Logical Instructions (1)

---

## Bitwise AND, OR, XOR

- Syntax

- and <reg>, <reg>
- and <mem>, <reg>
- and <reg>, <mem>
- and <imm>, <reg>
- and <imm>, <mem>
- Similar syntaxes for OR and XOR

- Examples

- and `$0x0F, %eax` /\* clear all but the last 4 bits of EAX. \*/
- xor `%edx, %edx` /\* set the contents of EDX to zero. \*/



# Logical Instructions (2)

---

## Bitwise Logical NOT, 2's Complement Negation

- Syntax

- not <reg>
- not <mem>
- neg <reg>
- neg <mem>

- Examples

- not %eax /\* flip all the bits of EAX \*/
- neg %eax /\* EAX is set to (- EAX) \*/



# Logical Instructions (3)

---

## Shift Left (SHL) and Right (SHR, SAR)

- Syntax

- `shl <imm8>, <reg>`
- `shl <imm8>, <mem>`
- `shl %cl, <reg>`
- `shl %cl, <mem>`
- Similar syntaxes for SHR and SAR

- Examples

- `shl $1, %eax /* EAX = EAX << 2  $\equiv$  EAX *= 2 (if the most significant bit is 0) */`
- `shr %cl, %ebx /* EBX = EBX >> CL (LOGICAL SHIFT TO RIGHT) */`
- `sar %cl, %ebx /* EBX = EBX >> CL  $\equiv$  EBX <- floor(EBX/(2^CL)) */`



# Control Flow Instructions (1)

---

Modify the Instruction Pointer (EIP) based on result of the last arithmetic state

- Syntax
  - `cmp <reg>, <reg>` (Equivalent to `sub <reg>, <reg>` but does not change the destination value)
  - `cmp <mem>, <reg>`
  - `cmp <reg>, <mem>`
  - `cmp <imm>, <reg>`
  - `je <label>` (jump when equal)
  - `jne <label>` (jump when not equal)
  - `jz <label>` (jump when last result was zero)
  - `jnz <label>` (jump when last result was non-zero)
  - `jg <label>` (jump when greater than)
  - `jge <label>` (jump when greater than or equal to)
  - `jl <label>` (jump when less than)
  - `jle <label>` (jump when less than or equal to)
  - `jmp <label>` (unconditional jump)



# Control Flow Instructions (2)

- Example - A count-up FOR loop:

**begin:**

```
xor %ecx, %ecx /* Zero out the counter register */  
mov (%esi), %eax /* Store the final count in EAX */
```

**loop:**

```
/* DO SOMETHING HERE */  
inc %ecx  
cmp %eax, %ecx  
jl loop /* if counter < final_count then jump to loop */
```

- Example - A count-down FOR loop:

**begin:**

```
mov (%esi), %ecx /* Store the count in ECX */
```

**1:**

```
/* DO SOMETHING HERE */  
dec %ecx  
jnz 1b /* if ECX != 0 then jump to label 1 before this line */
```



# Control Flow Instructions (3)

---

- Call/Ret : Call a function/Return from a function

- Syntax

- call <label>
  - ret

- call is equivalent to

```
push %eip /* Save the address of the next instruction on top of the stack */  
jmp foo   /* jump to the label foo, i.e., name of the destination function */
```

- ret is equivalent to

```
pop %eip /* Retrieve the return address from the stack-top */
```

- More on x86 Calling Convention: [Here](#)



# GNU Inline Assembly

---

- Can't do everything in C! E.g., How to make a system call (INT 0x80) in C? :/
- So, what to do now?
  1. Make object files from Assembly and link them with your C code!
  2. Use inline Assembly!



# GNU Inline Assembly

---

- Basic syntax: You can have something like the following in your C functions:

- `asm ("assembly instruction");`
  - `asm ("assembly instruction 1\n\t"  
"assembly instruction 2\n\t");`

- Example

```
asm ("pushl %eax\n\t"  
    "movl $0, %eax\n\t"  
    "popl %eax");
```



# GNU Inline Assembly

---

- Extended syntax: We can let the compiler decide what registers to use and transfer data between registers and our C variables!
  - `asm ("statements" : outputs : inputs : clobbered);`
- Format of outputs and inputs:
  - `"flags"(variable_name), "flags"(variable_name), ...`
- Flags:
  - `"r"` or `"q"`: Use a register for as an input operand
  - `"=r"` or `"=q"`: Use a register as an output operand
  - `"m"`, `"=m"`: Memory input/output operand
  - `"a"`, `"b"`, `"c"`, `"d"`, `"S"`, `"D"`, `"N"`: Use registers EAX, EBX, ECX, EDX, ESI, EDI, and 0-255 immediate value, respectively



# GNU Inline Assembly

---

- Example 1:  $x\_times\_5 = x + (4 * x)$ 
  - `asm ("leal (%1,%1,4), %0" : "=r"(x_times_5) : "r"(x) );`
  - `asm ("leal (%ebx,%%ebx,4), %%ebx" : "=b"(x) : "b"(x) );`
- Example 2: Read one byte from an I/O port
  - `asm ("inb %1,%0" : "=a"(value) : "Nd"(port_no));`
- More info? [Here](#)
- Still need more info? [Here](#)



# References

---

- [X86 Assembly Guide](#) by the Flint Group @ Yale
- [Brennan's Guide to Inline Assembly](#)
- [GCC-Inline-Assembly-HOWTO](#) by Sandeep.S